

Chapter 6: CPU Scheduling

Chapter 6: CPU Scheduling

COURSE INSTRUCTORS:

DR. MUHAMMAD NASEEM, ENGR. FARHEEN QAZI, MS. FALAK SALEEM

COURSE : OPERATING SYSTEMS (SWE-204)

BATCH : 2019 FALL : 2020

DEPARTMENT OF SOFTWARE ENGINEERING
SIR SYED UNIVERSITY OF ENGINEERING & TECHNOLOGY

1

Today's Agenda

- ▶ Basic Concepts
- ▶ Scheduling Criteria
- ▶ Scheduling Algorithms
- ▶ Multilevel Queue
- ▶ Multilevel Feedback Queue
- ▶ Algorithm Evaluation

2

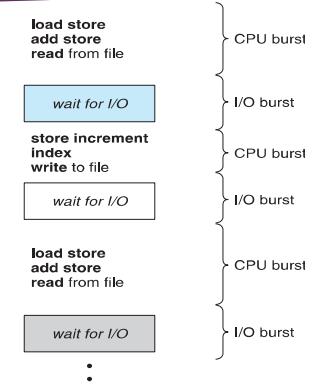
Objectives

- ▶ To introduce CPU scheduling, which is the basis for multiprogrammed operating systems
- ▶ To describe various CPU-scheduling algorithms
- ▶ To discuss evaluation criteria for selecting a CPU-scheduling algorithm for a particular system

3

Basic Concepts

- ▶ Maximum CPU utilization obtained with multiprogramming
- ▶ CPU-I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- ▶ **CPU burst** followed by **I/O burst**
- ▶ CPU burst distribution is of main concern



4

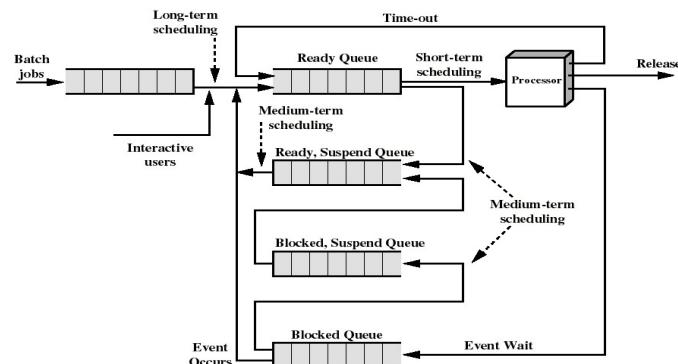
Chapter 6: CPU Scheduling

CPU Scheduler

- ▶ Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them.
- ▶ CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- ▶ Scheduling under 1 and 4 is **non-preemptive**
- ▶ All other scheduling is **preemptive**

5

Queuing Diagram for Scheduling



6

Dispatcher

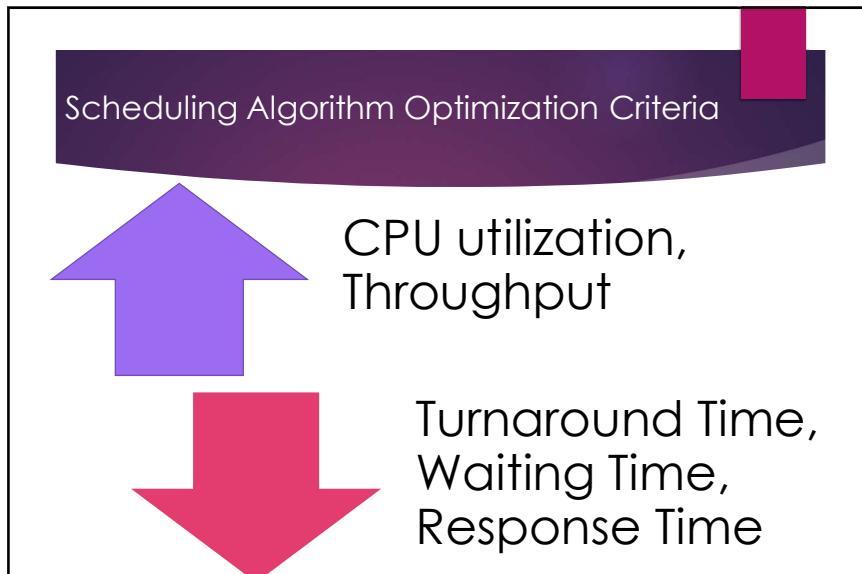
- ▶ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- ▶ **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running

7

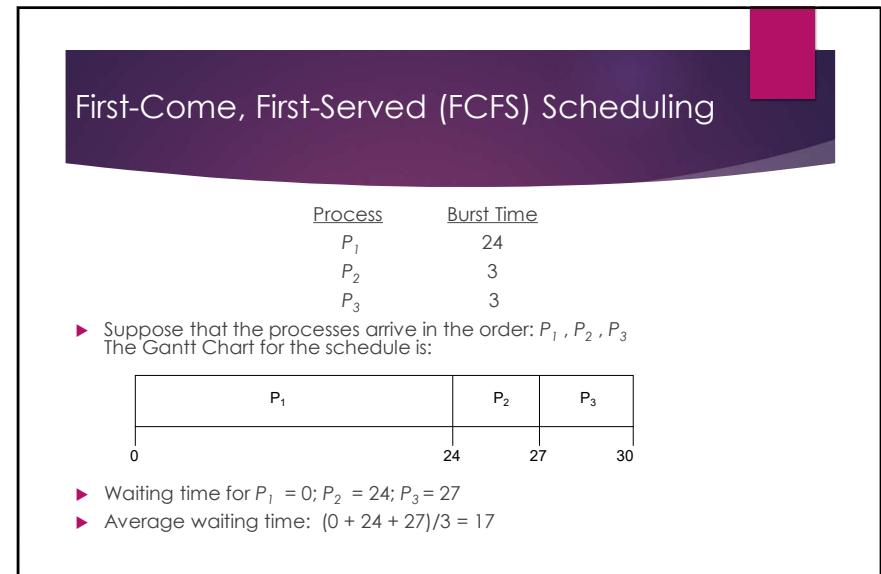
Scheduling Criteria

- ▶ **CPU utilization** – keep the CPU as busy as possible
- ▶ **Throughput** – # of processes that complete their execution per time unit
- ▶ **Turnaround time** – amount of time to execute a particular process
- ▶ **Waiting time** – amount of time a process has been waiting in the ready queue
- ▶ **Response time** – amount of time it takes from when a request was submitted until the first response is produced, not output (for time-sharing environment)

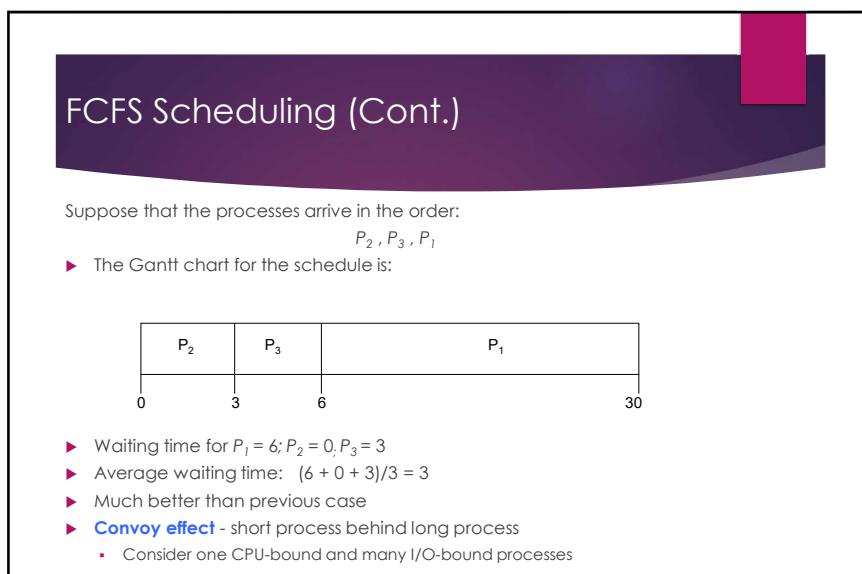
8



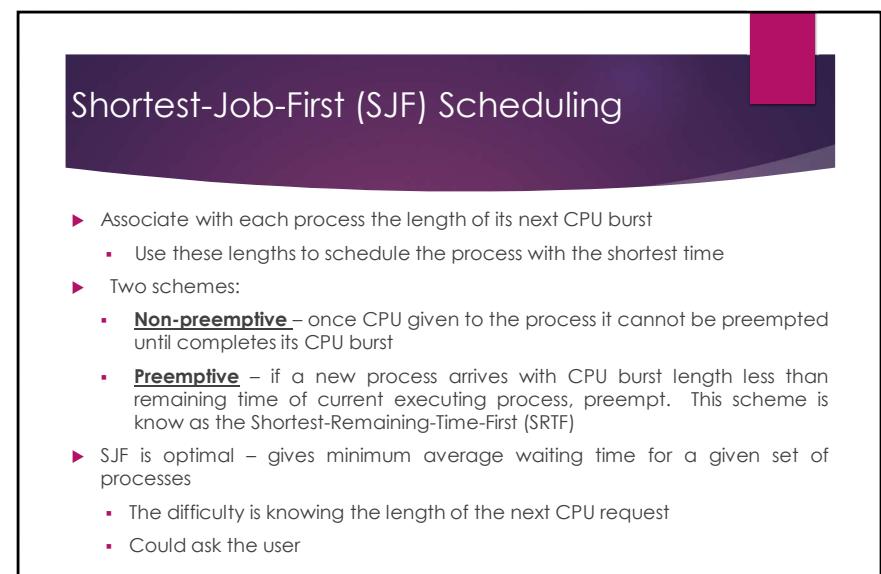
9



10



11

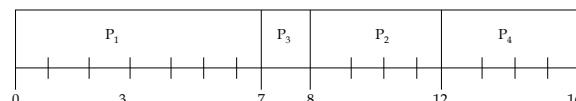


12

Example of SJF (non-preemptive)

Process	Arrival Time	Burst Time
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

► SJF (non-preemptive)



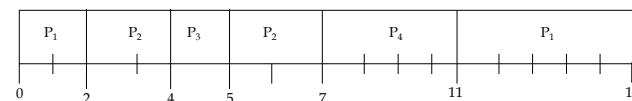
- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

13

Example of SJF (preemptive)

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

► SJF (preemptive)



► Average waiting time $\equiv (9 + 1 + 0 + 2)/4 \equiv 3$

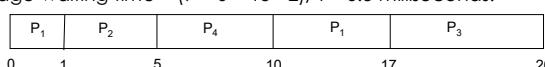
14

Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- Preemptive SJF Gantt Chart
 - Average waiting time = $(9 + 0 + 15 + 2)/4 = 6.5 \text{ milliseconds}$



- Average waiting time = $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$ msec

15

Priority Scheduling

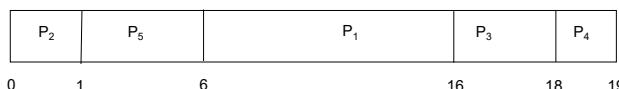
- ▶ A priority number (integer) is associated with each process
 - ▶ The CPU is allocated to the process with the highest priority (smallest integer = highest priority)
 - Preemptive
 - Non-preemptive
 - ▶ SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
 - ▶ Problem = **Starvation** – low priority processes may never execute
 - ▶ Solution = **Aging** – as time progresses increase the priority of the process

16

Example of Priority Scheduling (non-preemptive)

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

- ▶ Priority scheduling Gantt Chart



- ▶ Average waiting time = 8.2 msec

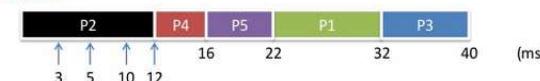
17

Example#02: Priority (non-preemptive)

❖ Example (Non-Preemptive)

Tasks	Arrival Times (ms)	Burst Time (ms)	Priority
P2	0.0	12	2
P3	3.0	8	5
P4	5.0	4	1
P1	10.0	10	4
P5	12.0	6	3

❖ Gantt chart



18

Contd....

❖ Time Calculation (Non-Preemptive)

$$P1 : 22 - 10 = 12 \text{ ms}$$

$$P2 : 0 \text{ ms}$$

$$P3 : 32 - 3 = 29 \text{ ms}$$

$$P4 : 12 - 5 = 7 \text{ ms}$$

$$P5 : 16 - 12 = 4 \text{ ms}$$

$$\text{Total waiting time : } P1 + P2 + P3 + P4 + P5 = 52 \text{ ms}$$

$$\text{Average waiting time : } 52 \text{ ms} / 5 = 10.4 \text{ ms}$$

19

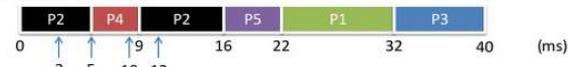
Example of Priority Scheduling (preemptive)

❖ Available in Preemptive and Non-Preemptive mode

❖ Example (Preemptive)

Tasks	Arrival Times (ms)	Burst Time (ms)	Priority
P2	0.0	12	2
P3	3.0	8	5
P4	5.0	4	1
P1	10.0	10	4
P5	12.0	6	3

❖ Gantt chart



20

Contd....

❖ Time Calculation (Non-Preemptive)

$$P_1 : 22 - 10 = 12 \text{ ms}$$

$$P_2 : 0 \text{ ms}$$

$$P_3 : 32 - 3 = 29 \text{ ms}$$

$$P_4 : 12 - 5 = 7 \text{ ms}$$

$$P_5 : 16 - 12 = 4 \text{ ms}$$

$$\text{Total waiting time} : P_1 + P_2 + P_3 + P_4 + P_5 = 52 \text{ ms}$$

$$\text{Average waiting time} : 52 \text{ ms} / 5 = 10.4 \text{ ms}$$

21

Round Robin (RR)

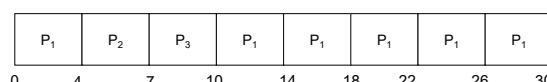
- ▶ Each process gets a small unit of CPU time (**time quantum** q), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ▶ If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once. No process waits more than $(n-1)q$ time units.
- ▶ Timer interrupts every quantum to schedule next process
- ▶ Performance
 - q large \Rightarrow FIFO
 - q small \Rightarrow q must be large with respect to context switch, otherwise overhead is too high

22

Example of RR with Time Quantum = 4

Process	Burst Time
P_1	24
P_2	3
P_3	3

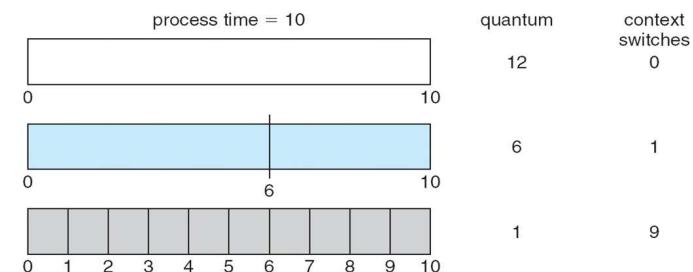
- ▶ The Gantt chart is:



- ▶ Typically, higher average turnaround than SJF, but better **response**
- ▶ q should be large compared to context switch time
- ▶ q usually 10ms to 100ms, context switch < 10 usec

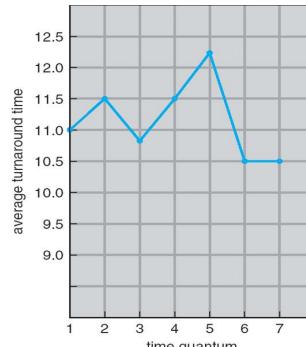
23

Time Quantum and Context Switch Time



24

Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

80% of CPU bursts should be shorter than q

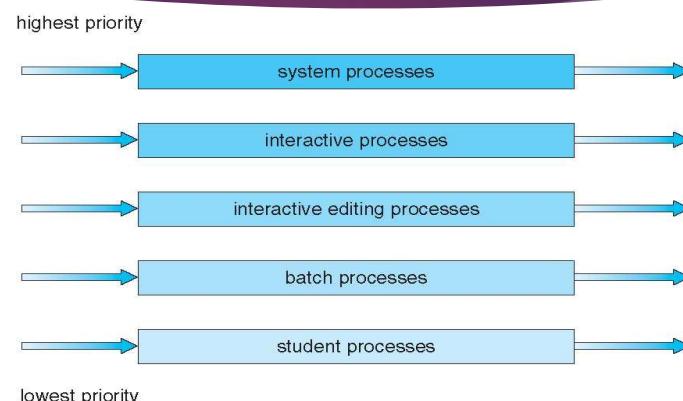
25

Multilevel Queue

- ▶ Ready queue is partitioned into separate queues, eg:
 - **foreground** (interactive)
 - **background** (batch)
- ▶ Process permanently in a given queue
- ▶ Each queue has its own scheduling algorithm:
 - foreground – RR
 - background – FCFS
- ▶ Scheduling must be done between the queues:
 - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - 20% to background in FCFS

26

Multilevel Queue Scheduling



27

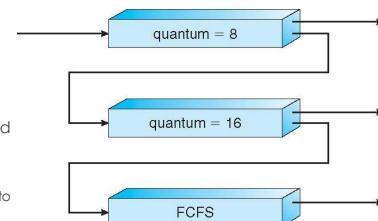
Multilevel Feedback Queue

- ▶ A process can move between the various queues; aging can be implemented this way
- ▶ Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

28

Example of Multilevel Feedback Queue

- ▶ Three queues:
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- ▶ Scheduling
 - A new job enters queue Q_0 which is served FCFS
 - When it gains CPU, job receives 8 milliseconds
 - If it does not finish in 8 milliseconds, job is moved to queue Q_1
 - At Q_1 , job is again served FCFS and receives 16 additional milliseconds
 - If it still does not complete, it is preempted and moved to queue Q_2



29

Algorithm Evaluation

- ▶ How to select CPU-scheduling algorithm for an OS?
- ▶ Determine criteria, then evaluate algorithms
- ▶ Deterministic modeling
 - Type of **analytic evaluation**
 - Takes a particular predetermined workload and defines the performance of each algorithm for that workload

30

Queuing Models

- ▶ Describes the arrival of processes, and CPU and I/O bursts probabilistically
 - Commonly exponential, and described by mean
 - Computes average throughput, utilization, waiting time, etc
- ▶ Computer system described as network of servers, each with queue of waiting processes
 - Knowing arrival rates and service rates
 - Computes utilization, average queue length, average wait time, etc

31

Class Activity # 01

Draw the Gantt Chart and also calculate average waiting time and average turnaround time of the following algorithms using given table:

1. FCFS
2. SJF(Non-Preemptive)
3. SJF(Preemptive)
4. Round Robin (Quantum = 5).

Process	Arrival Time	Burst Time
P1	0.0	16
P2	0.3	5
P3	1.0	3

32

Class Activity # 02

Draw the Gantt Chart and also calculate average waiting time and average turnaround time of the following algorithms using given table:

1. FCFS
2. SJF(Non-Preemptive)
3. SJF(Preemptive)
4. Priority(Non-Preemptive)
5. Priority(Preemptive)
6. Round Robin (Quantum = 4).

Process	Arrival Time	Burst Time	Priority
P1	2	6	4
P2	5	3	1
P3	1	8	2
P4	0	3	1.5
P5	4	4	3.5

33

Class Activity # 03

Draw the Gantt Chart and also calculate average waiting time and average turnaround time of the following algorithms using given table:

1. FCFS
2. SJF(Non-Preemptive)
3. SJF(Preemptive)
4. Priority(Non-Preemptive)
5. Priority(Preemptive)
6. Round Robin (Quantum = 4).

Process	Arrival Time	Burst Time	Priority
P1	1	5	1
P2	7	3	2
P3	9	4	7
P4	12	2	4
P5	16	10	1.5

34

End of Chapter 6

Thank you

35